

credit_risk_ensemble

September 24, 2019

```
[1]: import warnings
warnings.filterwarnings('ignore')

[2]: import numpy as np
import pandas as pd
from pathlib import Path
from collections import Counter

[3]: from sklearn.metrics import balanced_accuracy_score
from sklearn.metrics import confusion_matrix
from imblearn.metrics import classification_report_imbalanced
```

1 Read the CSV and Perform Basic Data Cleaning

```
[4]: # https://help.lendingclub.com/hc/en-us/articles/215488038-What-do-the-different-Note-statuses-mean-

columns = [
    "loan_amnt", "int_rate", "installment", "home_ownership",
    "annual_inc", "verification_status", "issue_d", "loan_status",
    "pymnt_plan", "dti", "delinq_2yrs", "inq_last_6mths",
    "open_acc", "pub_rec", "revol_bal", "total_acc",
    "initial_list_status", "out_prncp", "out_prncp_inv", "total_pymnt",
    "total_pymnt_inv", "total_rec_prncp", "total_rec_int", "total_rec_late_fee",
    "recoveries", "collection_recovery_fee", "last_pymnt_amnt", "next_pymnt_d",
    "collections_12_mths_ex_med", "policy_code", "application_type",
    "acc_now_delinq",
    "tot_coll_amt", "tot_cur_bal", "open_acc_6m", "open_act_il",
    "open_il_12m", "open_il_24m", "mths_since_rcnt_il", "total_bal_il",
    "il_util", "open_rv_12m", "open_rv_24m", "max_bal_bc",
    "all_util", "total_rev_hi_lim", "inq_fi", "total_cu_tl",
    "inq_last_12m", "acc_open_past_24mths", "avg_cur_bal", "bc_open_to_buy",
    "bc_util", "chargeoff_within_12_mths", "delinq_amnt", "mo_sin_old_il_acct",
    "mo_sin_old_rev_tl_op", "mo_sin_rcnt_rev_tl_op", "mo_sin_rcnt_tl",
    "mort_acc",
```

```

    "mths_since_recent_bc", "mths_since_recent_inq", "num_accts_ever_120_pd",
    → "num_actv_bc_tl",
    "num_actv_rev_tl", "num_bc_sats", "num_bc_tl", "num_il_tl",
    "num_op_rev_tl", "num_rev_accts", "num_rev_tl_bal_gt_0",
    "num_sats", "num_tl_120dpd_2m", "num_tl_30dpd", "num_tl_90g_dpd_24m",
    "num_tl_op_past_12m", "pct_tl_nvr_dlq", "percent_bc_gt_75",
    → "pub_rec_bankruptcies",
    "tax_liens", "tot_hi_cred_lim", "total_bal_ex_mort", "total_bc_limit",
    "total_il_high_credit_limit", "hardship_flag", "debt_settlement_flag"
]

```

```
target = ["loan_status"]
```

```

[5]: # Load the data
file_path = Path('../Resources/LoanStats_2019Q1.csv.zip')
df = pd.read_csv(file_path, skiprows=1)[-2]
df = df.loc[:, columns].copy()

# Drop the null columns where all values are null
df = df.dropna(axis='columns', how='all')

# Drop the null rows
df = df.dropna()

# Remove the `Issued` loan status
issued_mask = df['loan_status'] != 'Issued'
df = df.loc[issued_mask]

# convert interest rate to numerical
df['int_rate'] = df['int_rate'].str.replace('%', '')
df['int_rate'] = df['int_rate'].astype('float') / 100

# Convert the target column values to low_risk and high_risk based on their
→ values
x = {'Current': 'low_risk'}
df = df.replace(x)

x = dict.fromkeys(['Late (31-120 days)', 'Late (16-30 days)', 'Default', 'In_
→ Grace Period'], 'high_risk')
df = df.replace(x)

df.reset_index(inplace=True, drop=True)

df.head()

```

```

[5]:  loan_amnt  int_rate  installment  home_ownership  annual_inc  \
0      10500.0    0.1719         375.35             RENT      66000.0
1      25000.0    0.2000         929.09             MORTGAGE    105000.0
2      20000.0    0.2000         529.88             MORTGAGE    56000.0
3      10000.0    0.1640         353.55             RENT      92000.0
4      22000.0    0.1474         520.39             MORTGAGE    52000.0

      verification_status  issue_d  loan_status  pymnt_plan    dti  ...  \
0      Source Verified  Mar-2019    low_risk          n  27.24  ...
1      Verified  Mar-2019    low_risk          n  20.23  ...
2      Verified  Mar-2019    low_risk          n  24.26  ...
3      Verified  Mar-2019    low_risk          n  31.44  ...
4      Not Verified  Mar-2019    low_risk          n  18.76  ...

      pct_tl_nvr_dlq  percent_bc_gt_75  pub_rec_bankruptcies  tax_liens  \
0           85.7           100.0           0.0           0.0
1           91.2           50.0           1.0           0.0
2           66.7           50.0           0.0           0.0
3          100.0           50.0           1.0           0.0
4          100.0           0.0           0.0           0.0

      tot_hi_cred_lim  total_bal_ex_mort  total_bc_limit  \
0          65687.0          38199.0          2000.0
1          271427.0          60641.0          41200.0
2          60644.0          45684.0          7500.0
3          99506.0          68784.0          19700.0
4          219750.0          25919.0          27600.0

      total_il_high_credit_limit  hardship_flag  debt_settlement_flag
0           61987.0              N              N
1           49197.0              N              N
2           43144.0              N              N
3           76506.0              N              N
4           20000.0              N              N

```

[5 rows x 86 columns]

2 Split the Data into Training and Testing

```

[6]: # Create our features
      X = # YOUR CODE HERE

      # Create our target
      y = # YOUR CODE HERE

```

```

[7]: X.describe()

```

```
[7]:      loan_amnt      int_rate      installment      annual_inc      dti \
count  68817.000000  68817.000000  68817.000000  6.881700e+04  68817.000000
mean   16677.594562    0.127718    480.652863  8.821371e+04   21.778153
std    10277.348590    0.048130    288.062432  1.155800e+05   20.199244
min     1000.000000    0.060000     30.890000  4.000000e+01    0.000000
25%     9000.000000    0.088100    265.730000  5.000000e+04   13.890000
50%    15000.000000    0.118000    404.560000  7.300000e+04   19.760000
75%    24000.000000    0.155700    648.100000  1.040000e+05   26.660000
max    40000.000000    0.308400   1676.230000  8.797500e+06  999.000000
```

```
      delinq_2yrs  inq_last_6mths      open_acc      pub_rec \
count  68817.000000  68817.000000  68817.000000  68817.000000
mean     0.217766     0.497697    12.587340    0.126030
std     0.718367     0.758122     6.022869    0.336797
min     0.000000     0.000000     2.000000    0.000000
25%     0.000000     0.000000     8.000000    0.000000
50%     0.000000     0.000000    11.000000    0.000000
75%     0.000000     1.000000    16.000000    0.000000
max     18.000000     5.000000    72.000000    4.000000
```

```
      revol_bal  ...  issue_d_Mar-2019  pymnt_plan_n \
count  68817.000000  ...    68817.000000    68817.0
mean   17604.142828  ...         0.177238         1.0
std    21835.880400  ...         0.381873         0.0
min      0.000000  ...         0.000000         1.0
25%     6293.000000  ...         0.000000         1.0
50%    12068.000000  ...         0.000000         1.0
75%    21735.000000  ...         0.000000         1.0
max   587191.000000  ...         1.000000         1.0
```

```
      initial_list_status_f  initial_list_status_w  next_pymnt_d_Apr-2019 \
count      68817.000000      68817.000000      68817.000000
mean         0.123879         0.876121         0.383161
std         0.329446         0.329446         0.486161
min         0.000000         0.000000         0.000000
25%         0.000000         1.000000         0.000000
50%         0.000000         1.000000         0.000000
75%         0.000000         1.000000         1.000000
max         1.000000         1.000000         1.000000
```

```
      next_pymnt_d_May-2019  application_type_Individual \
count      68817.000000      68817.000000
mean         0.616839         0.860340
std         0.486161         0.346637
min         0.000000         0.000000
25%         0.000000         1.000000
50%         1.000000         1.000000
```

75%	1.000000	1.000000
max	1.000000	1.000000

	application_type_Joint App	hardship_flag_N	debt_settlement_flag_N
count	68817.000000	68817.0	68817.0
mean	0.139660	1.0	1.0
std	0.346637	0.0	0.0
min	0.000000	1.0	1.0
25%	0.000000	1.0	1.0
50%	0.000000	1.0	1.0
75%	0.000000	1.0	1.0
max	1.000000	1.0	1.0

[8 rows x 95 columns]

```
[8]: # Check the balance of our target values
y['loan_status'].value_counts()
```

```
[8]: low_risk      68470
high_risk       347
Name: loan_status, dtype: int64
```

```
[9]: # Split the X and y into X_train, X_test, y_train, y_test
# YOUR CODE HERE
```

3 Ensemble Learners

In this section, you will compare two ensemble algorithms to determine which algorithm results in the best performance. You will train a Balanced Random Forest Classifier and an Easy Ensemble AdaBoost classifier . For each algorithm, be sure to complete the following steps:

1. Train the model using the training data.
2. Calculate the balanced accuracy score from sklearn.metrics.
3. Print the confusion matrix from sklearn.metrics.
4. Generate a classification report using the imbalanced_classification_report from imbalanced-learn.
5. For the Balanced Random Forest Classifier onely, print the feature importance sorted in descending order (most important feature to least important) along with the feature score

Note: Use a random state of 1 for each algorithm to ensure consistency between tests

3.0.1 Balanced Random Forest Classifier

```
[10]: # Resample the training data with the RandomOversampler
# YOUR CODE HERE
```

```
[10]: BalancedRandomForestClassifier(bootstrap=True, class_weight=None,
criterion='gini', max_depth=None, max_features='auto',
max_leaf_nodes=None, min_impurity_decrease=0.0,
```

```

min_samples_leaf=2, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=1,
oob_score=False, random_state=1, replacement=False,
sampling_strategy='auto', verbose=0, warm_start=False)

```

```

[11]: # Calculated the balanced accuracy score
      # YOUR CODE HERE

```

```

[11]: 0.7855052723466922

```

```

[12]: # Display the confusion matrix
      # YOUR CODE HERE

```

```

[12]: array([[ 68,   33],
            [1749, 15355]])

```

```

[13]: # Print the imbalanced classification report
      # YOUR CODE HERE

```

	pre	rec	spe	f1	geo	iba
sup						
high_risk	0.04	0.67	0.90	0.07	0.78	0.59
101						
low_risk	1.00	0.90	0.67	0.95	0.78	0.62
17104						
avg / total	0.99	0.90	0.67	0.94	0.78	0.62
17205						

```

[14]: # List the features sorted in descending order by feature importance
      # YOUR CODE HERE

```

```

loan_amnt: (0.09175752102205247)
int_rate: (0.06410003199501778)
installment: (0.05764917485461809)
annual_inc: (0.05729679526683975)
dti: (0.05174788106507317)
delinq_2yrs: (0.031955619175665397)
inq_last_6mths: (0.02353678623968216)
open_acc: (0.017078915518993903)
pub_rec: (0.017014861224701222)
revol_bal: (0.016537957646730293)
total_acc: (0.016169718411077325)
out_prncp: (0.01607049983545137)
out_prncp_inv: (0.01599866290723441)
total_pymnt: (0.015775537221600675)
total_pymnt_inv: (0.01535560674178928)

```

total_rec_prncp: (0.015029265003541079)
total_rec_int: (0.014828006488636946)
total_rec_late_fee: (0.01464881608833323)
recoveries: (0.014402430445752665)
collection_recovery_fee: (0.014318832248876989)
last_pymnt_amnt: (0.013519867193755364)
collections_12_mths_ex_med: (0.013151520216882331)
policy_code: (0.013101578263049833)
acc_now_delinq: (0.012784600558682344)
tot_coll_amt: (0.012636608914961465)
tot_cur_bal: (0.012633464965390648)
open_acc_6m: (0.012406321468566728)
open_act_il: (0.011687404692448701)
open_il_12m: (0.01156494245653799)
open_il_24m: (0.011455878011762288)
mths_since_rcnt_il: (0.011409157520644688)
total_bal_il: (0.01073641504525053)
il_util: (0.010380085181706624)
open_rv_12m: (0.010097528131347774)
open_rv_24m: (0.00995373830638152)
max_bal_bc: (0.00991410213601043)
all_util: (0.009821715826953788)
total_rev_hi_lim: (0.009603648248133598)
inq_fi: (0.009537423049553)
total_cu_tl: (0.008976776055926955)
inq_last_12m: (0.008870623013604539)
acc_open_past_24mths: (0.008745106187024114)
avg_cur_bal: (0.008045578273709669)
bc_open_to_buy: (0.007906251501807723)
bc_util: (0.00782073260901301)
chargeoff_within_12_mths: (0.007798696767389274)
delinq_amnt: (0.007608045628523077)
mo_sin_old_il_acct: (0.0075861537897335815)
mo_sin_old_rev_tl_op: (0.007554511001273182)
mo_sin_rcnt_rev_tl_op: (0.007471884930172615)
mo_sin_rcnt_tl: (0.007273779915807858)
mort_acc: (0.006874845464745796)
mths_since_recent_bc: (0.006862142977394886)
mths_since_recent_inq: (0.006838718858820505)
num_accts_ever_120_pd: (0.006413554699909871)
num_actv_bc_tl: (0.006319439816216779)
num_actv_rev_tl: (0.006160469432535709)
num_bc_sats: (0.006066257227997291)
num_bc_tl: (0.005981472544437747)
num_il_tl: (0.0055301594524349495)
num_op_rev_tl: (0.004961823663836347)
num_rev_accts: (0.004685198497435334)
num_rev_tl_bal_gt_0: (0.0045872929977180356)

```

num_sats: (0.0041651633321967895)
num_tl_120dpd_2m: (0.004016461341161775)
num_tl_30dpd: (0.0032750717701661657)
num_tl_90g_dpd_24m: (0.0027565184136781346)
num_tl_op_past_12m: (0.0026174030074401656)
pct_tl_nvr_dlq: (0.002279671873697176)
percent_bc_gt_75: (0.0021899772867773103)
pub_rec_bankruptcies: (0.0020851101815353096)
tax_liens: (0.0018404849590376573)
tot_hi_cred_lim: (0.001736019018028134)
total_bal_ex_mort: (0.0015472230884974506)
total_bc_limit: (0.0012263315437383057)
total_il_high_credit_limit: (0.0012213148580230454)
home_ownership_ANY: (0.0012151288883862276)
home_ownership_MORTGAGE: (0.0008976722260399365)
home_ownership_OWN: (0.0008125182396705508)
home_ownership_RENT: (0.000573414997420326)
verification_status_Not Verified: (0.0005168345750594915)
verification_status_Source Verified: (0.0004192455022893127)
verification_status_Verified: (0.0)
issue_d_Feb-2019: (0.0)
issue_d_Jan-2019: (0.0)
issue_d_Mar-2019: (0.0)
pymnt_plan_n: (0.0)
initial_list_status_f: (0.0)
initial_list_status_w: (0.0)
next_pymnt_d_Apr-2019: (0.0)
next_pymnt_d_May-2019: (0.0)
application_type_Individual: (0.0)
application_type_Joint App: (0.0)
hardship_flag_N: (0.0)
debt_settlement_flag_N: (0.0)

```

3.0.2 Easy Ensemble AdaBoost Classifier

```

[15]: # Train the Classifier
      # YOUR CODE HERE

```

```

[15]: EasyEnsembleClassifier(base_estimator=None, n_estimators=100, n_jobs=1,
                             random_state=1, replacement=False, sampling_strategy='auto',
                             verbose=0, warm_start=False)

```

```

[16]: # Calculated the balanced accuracy score
      # YOUR CODE HERE

```

```

[16]: 0.9316600714093861

```

```

[17]: # Display the confusion matrix
      # YOUR CODE HERE

```



```
[17]: array([[ 93,    8],
           [ 983, 16121]])
```

```
[18]: # Print the imbalanced classification report
      # YOUR CODE HERE
```

	pre	rec	spe	f1	geo	iba
sup						
high_risk	0.09	0.92	0.94	0.16	0.93	0.87
101						
low_risk	1.00	0.94	0.92	0.97	0.93	0.87
17104						
avg / total	0.99	0.94	0.92	0.97	0.93	0.87
17205						

```
[ ]:
```